

Lecture 1: Linear Proofs and Quadratic Proofs

Samuel B. Hopkins

August 31, 2024




Figure 1: This course in a nutshell.

Reader beware: these notes have not been subjected to the usual scrutiny reserved for peer-reviewed publication, and may contain errors.

1 Minimum and Maximum Cuts

Let's start with the minimum and maximum cut problems.

Problem 1 (Min- (s, t) -Cut). *Given a graph $G = (V, E)$ and $s, t \in V$, find $S \subseteq V$ such that $s \in S$ and $t \in \bar{S}$ minimizing the number of cut edges $E(S, \bar{S})$.*

Min-cut can be solved in polynomial time, e.g. by using Ford-Fulkerson or another max-flow algorithm to find the maximum flow in the graph from s, t , then finding all the vertices reachable from s in this flow's residual graph.

Problem 2 (Max-Cut). *Given a graph $G = (V, E)$, find $S \subseteq V$ maximizing the number of cut edges $E(S, \bar{S})$.*

Unlike its friend Min-Cut, Max-Cut is NP-hard, so we do not expect to design a polynomial-time algorithm for it. Instead, we will aim for an *approximation algorithm*. As a reminder, an α -approximation algorithm ($\alpha \geq 1$) for Max-Cut takes a graph G and outputs a cut $S \subseteq V$ such that $E(S, \bar{S}) \geq \frac{1}{\alpha} \cdot \text{OPT}$, where OPT is the number of cut edges in the optimal solution.

2-approximation for Max-Cut The simplest approximation algorithm for Max-Cut is the following: include each $v \in V$ in S independently with probability $1/2$.

It is easy to show that every edge (i, j) is cut with probability $1/2$, so the expected number of cut edges is $|E|/2$. At the same time, the optimum value $\text{OPT} \leq |E|$ – we can't do better than cutting all of the edges. So

$$\mathbb{E}_S E(S, \bar{S}) = \frac{|E|}{2} \geq \frac{\text{OPT}}{2},$$

and hence the algorithm is a 2-approximation. As an aside, this algorithm can be derandomized into a greedy algorithm.

Question. *Can we do better than a 2-approximation for Max-Cut?*

To investigate, we will take a detour to the world of *proofs* and *proof systems*.

2 Proofs and (Approximation) Algorithms

Every approximation algorithm is a proof system for proving bounds on OPT.

Let us explain what is meant by this sentence in the contexts of Min-Cut and Max-Cut, staying informal for now about what we mean by a “proof system.”

2.1 Min-Cut

Upper bounds on Min-Cut Suppose G is a graph held by a prover, Alice, who claims that $\text{Min-Cut}(G) \leq k$. Alice can convince a polynomial-time verifier, Bob, of this fact, by supplying him with a cut $S \subseteq V$ such that $E(S, \bar{S}) \leq k$. Given the cut S , Bob can verify that $E(S, \bar{S}) \leq k$ in polynomial time. This is the argument which shows Min-Cut is in NP.

Lower bounds on Min-Cut What if Alice claims that $\text{Min-Cut}(G) \geq k$? What kind of proof could she show Bob? As you might recall from undergraduate algorithms, one type of proof Alice could use is a valid flow from s to t . If Alice can show Bob a flow of value k from s to t , meaning that she tells Bob how much flow to route on each edge of G , Bob can easily check that the flow is valid and that its value is k , and hence that $\text{Min-Cut}(G) \geq k$.

The max-flow min-cut theorem says that this kind of proof actually always exists – there is always an (s, t) flow of value $\text{Min-Cut}(G)$, and hence Alice can always convince Bob of any true statement of the form $\text{Min-Cut}(G) \geq k$ by showing him a flow of value k .

2.2 Max-Cut

Lower bounds on Max-Cut Similar to the case of upper bounds on the Min-Cut, suppose Alice claims that $\text{Max-Cut}(G) \geq k$. Alice can convince a polynomial-time verifier, Bob, of this fact, by simply supplying him with a cut $S \subseteq V$ such that $E(S, \bar{S}) \geq k$.

Upper bounds on OPT But now what if Alice claims that $\text{OPT} \leq k$? What kind of proof could she show Bob to convince him of this fact? Not so clear!

2.3 Approximation Algorithms are Proof Systems

Before we turn to particular kinds of proofs Alice might use, let’s justify the point of this detour: if there is a (deterministic) α -approximation algorithm ALG for Max-Cut, then there is a way for Alice to convince Bob that $\text{OPT} \leq \alpha \cdot k$ whenever $\text{OPT} \leq k$.¹

¹Or, in complexity-theoretic terms, $P \subseteq NP \cap coNP$.

One way would be for Bob to just run ALG – whatever cut S ALG outputs will have value $E(S, \bar{S}) \leq \text{OPT} \leq k$, and since ALG is an α -approximation algorithm, seeing ALG output S will convince Bob that $\text{OPT} \leq \alpha E(S, \bar{S}) \leq \alpha k$.

If we want to be more explicit about the proof system, we can imagine that Alice sends Bob a transcript of the execution of ALG, and Bob can verify that the transcript is correct and that the cut S output by ALG has value at most k , and hence that $\text{OPT} \leq \alpha k$.

So, if we design, say, a 1.99-approximation algorithm for Max-Cut, in the course of doing so we will have implicitly designed a proof system which Alice could use to convince Bob that $\text{OPT} \leq 1.99k$ whenever $\text{OPT} \leq k$. We might as well try to design the proof system explicitly.

3 Linear Proofs

Our goal is to design a proof system which Alice can use to convince Bob that $\text{Max-Cut}(G) \leq k$. To get inspiration, let's return to Min-Cut and have a more explicit look at the proof system for lower bounds on Min-Cut, through the lens of linear programming.

3.1 Linear Proofs and Min-Cut

We'll start by writing down a linear program for Min- (s, t) -Cut. Our LP will have variables x_v for $v \in V$, and also y_{uv} for edges $(u, v) \in E$. The intended meaning of the variables will be that $x_v = 1$ if $v \in S$ and $x_v = 0$ otherwise, and $y_{uv} = 1$ if (u, v) is a cut edge and $y_{uv} = 0$ otherwise. The LP is:

$$\min_{x, y} \sum_{(u, v) \in E} y_{uv} \text{ such that} \quad (1)$$

$$x_s = 1 \quad (2)$$

$$x_t = 0 \quad (3)$$

$$0 \leq y_{uv} \quad \forall (u, v) \in E \quad (4)$$

$$y_{uv} \geq x_u - x_v \quad \forall (u, v) \in E. \quad (5)$$

We are taking the tuples (u, v) to be ordered, and when there is an edge between u and v we assume E contains (u, v) and (v, u) . Notice that the final set of constraints encode the fact that $y_{uv} = 1$ only if (u, v) is a cut edge with $x_u = 1$ and $x_v = 0$. We won't prove it, but this LP is actually tight: the optimal value is exactly equal to the Min- (s, t) -Cut value.

The LP offers a proof system for proving lower bounds on the Min- (s, t) -Cut value via duality. As a reminder, a dual solution to the above LP consists of numbers $\lambda_s, \lambda_t \in \mathbb{R}$ (multipliers for the constraints $x_s = 1, x_t = 0$) and $\alpha_{u,v}, f_{uv} \geq 0$ for $(u, v) \in E$ (multipliers for the inequality constraints) such that

$$\sum_{(u, v) \in E} y_{uv} - k = \lambda_s(x_s - 1) + \lambda_t x_t + \sum_{(u, v) \in E} \alpha_{uv} y_{uv} + f_{uv}(y_{uv} - (x_u - x_v))$$

where here we mean equality *as linear functions of x and y* . Because the RHS of the inequality above must be nonnegative for every feasible solution x, y , the LHS must be as well, meaning that the Min- (s, t) -Cut value is at least k . If Alice shows Bob the dual solution, he can verify that it satisfies the above equality by checking that the coefficient of each variable x_i is the same on both sides.

This LP-based proof system is actually *equivalent* to the max-flow proof system from above! It's not too hard to prove that in order to satisfy the above equality, the dual variables f_{uv} have to be a valid t, s flow of value k , where we route f_{uv} units of flow from u to v .

Proof sketch that f_{uv} is a valid flow of value k . The coefficient of y_{uv} on the RHS is $\alpha_{uv} + f_{uv}$ and on the LHS it's 1. So we must have $1 = \alpha_{uv} + f_{uv}$, and hence $f_{uv} \in [0, 1]$. Thus f respects the edge capacities.

For $v \neq s, t$, the coefficient of x_v on the RHS is $\sum_{u \in N(v)} f_{uv} - f_{vu}$, and on the LHS it's 0. So we must have $\sum_{u \in N(v)} f_{uv} = \sum_{u \in N(v)} f_{vu}$ – that is, the incoming flow equals the outgoing flow at vertex v .

The constant term on the RHS is $-\lambda_s$ and on the LHS it's $-k$, so we must have $\lambda_s = k$. But, the coefficient of x_s on the RHS is $\lambda_s + \sum_{u \in N(s)} f_{su} - f_{us}$ and on the LHS it's 0. So we must have $\sum_{u \in N(s)} f_{su} - f_{us} = -k$, meaning that k units of flow are routed into s . \square

Furthermore, reformulating what we said before about the max-flow proof system, by the max-flow min-cut theorem, whenever $\text{Min-Cut}(G) \geq k$, Alice can prove this to Bob by exhibiting such λ, α , and f .

3.2 Linear Proofs and Max-Cut

A linear proof system turned out to be as powerful as we could hope for for Min-Cut. How about for Max-Cut? Of course we don't expect to come up with an exact LP formulation of Max-Cut, but perhaps we can arrive at a useful LP relaxation and a corresponding proof system nonetheless. (Even more optimistically, we might then hope to use that LP to go beyond a proof system and actually design an approximation algorithm.)

Unfortunately, this approach won't work out – in fact, one can show that there is no polynomial-size LP relaxation of Max-Cut which gives a $2 - \varepsilon$ approximation for any $\varepsilon > 0$ (although that result is beyond the scope of this class). But let's try to see why we can't just use the same type of LP as for Min-Cut.

As for Min-Cut, we can have variables x_v for $v \in V$ and y_{uv} for $(u, v) \in E$, with the intended meanings that $y_{uv} = 1$ if (u, v) is a cut edge and $y_{uv} = 0$ otherwise, and $x_u = 1$ if $u \in S$ and $x_u = 0$ otherwise. The objective function is easy enough to write: $\max \sum_{(u,v) \in E} y_{uv}$. But how can we constrain y_{uv} so that it is 1 only if (u, v) is a cut edge? In the Min-Cut case we used the fact that the LP was trying to minimize y_{uv} , together with the constraints $y_{uv} \geq x_u - x_v$, to force y_{uv} to act like the indicator for a cut edge. Now the LP is trying to maximize y_{uv} , so we can't use the same trick.

Of course, we could write down a silly LP relaxation of Max-Cut:

$$\max_y \sum_{(u,v) \in E} y_{uv} \text{ s.t. } 0 \leq y_{uv} \leq 1.$$

(Now we think of the (u, v) pairs as un-ordered.) The optimum value will always be $|E|$, meaning that this LP gives just a 2-approximation to the value of the Max-Cut. Just for fun, let's see the proof system that this LP gives us: there are dual variables $\alpha_{uv}, \lambda_{uv} \geq 0$, and to prove that $\text{Max-Cut}(G) \leq k$ we require that

$$k - \sum_{(u,v) \in E} y_{uv} = \sum_{(u,v) \in E} \alpha_{uv} y_{uv} + \lambda_{uv} (1 - y_{uv}).$$

Unfortunately, this linear proof system for Max-Cut will never let Alice prove anything better than $\text{Max-Cut}(G) \leq |E|$.

4 Quadratic Proofs

The linear proof system for Max-Cut was not very useful. Our LP relaxation didn't work well because we didn't have a good way to express the constraint that $y_{uv} = 1$ only if (u, v) is a cut edge.

Observe, however, that if we were not so tied to linear functions of our variables, we would not have this issue. It is traditional at this point to switch to placing the x_v variables in $\{-1, 1\}$. For $x_u, x_v \in \{-1, 1\}$ we could directly write the indicator $1(x_u \neq x_v)$ as a quadratic function of x_u, x_v :

$$1(x_u \neq x_v) = \frac{1}{2}(1 - x_u x_v)$$

so that the Max-Cut objective function would be

$$\max_x \frac{1}{2} \sum_{(u,v) \in E} \frac{1}{2} - \frac{x_u x_v}{2}.$$

Now let's see how we could use this quadratic objective function to define a proof system for Max-Cut. Suppose, as usual, that Alice wants to prove to Bob that $\text{Max-Cut}(G) \leq k$. In the linear proofs setting, Alice showed Bob a way to write the difference of k and the (linear) objective function as a sum of pieces which Bob could easily verify were nonnegative. The "pieces" were linear functions induced by constraints from the LP, which Bob knows must be nonnegative for any feasible solution y . It is important that Bob has such an *a priori* reason to believe that the pieces are nonnegative, because he does not have time to check nonnegativity for every feasible y .

Now the objective function is quadratic in x . What are some *quadratic* functions of x that Bob can easily check are nonnegative for any feasible solution x ?

Big idea: *The square of every linear function with real coefficients is everywhere nonnegative!*

So, suppose Alice could find m linear functions $f_1(x), \dots, f_m(x)$ – that is, each f_i has the form

$$f_i(x) = \sum_{v \in V} f_{iv} x_v$$

where $f_{iv} \in \mathbb{R}$ – such that for every $x \in \{-1, 1\}^V$,

$$k - \sum_{(u,v) \in E} \frac{1}{2} - \frac{x_u x_v}{2} = \sum_{i=1}^m f_i(x)^2.$$

Then, assuming that Bob could actually check that this equation holds for all $x \in \{-1, 1\}^V$, he would be convinced that $\text{Max-Cut}(G) \leq k$.

How could Bob convince himself that the equation holds for all $x \in \{-1, 1\}^V$, without spending time $2^{|V|}$? If Bob were trying to verify that the equation holds for every $x \in \mathbb{R}^V$, he could just check that the coefficient of every degree-2 monomial $x_u x_v$ on the LHS is equal to the coefficient of the same monomial on the RHS. It turns out that a similar approach works for $x \in \{-1, 1\}^V$, because of the multivariate discrete Fourier transform:

Fact 3. *Every function $f : \{-1, 1\}^n \rightarrow \mathbb{R}$ can be written uniquely as a multilinear polynomial in the x_i 's; that is, $f(x) = \sum_{S \subseteq [n]} \hat{f}_S \cdot \prod_{i \in S} x_i$ where $\hat{f}_S \in \mathbb{R}$.*

So, given f_1, \dots, f_m , Bob can expand $\sum_{i=1}^m f_i(x)^2$ in the monomial basis and replace each non-multilinear term x_v^2 with 1. Then he can compare the coefficients of each monomial on the LHS and RHS to verify that the equation holds for all $x \in \{-1, 1\}^V$. This verification procedure clearly runs in time polynomial in $|V|, m$.

We are left with (at least) three questions:

1. (Proof existence) For which k can Alice find a proof that $\text{Max-Cut}(G) \leq k$? In particular, do such proofs exist for $k < |E|$?
2. (Proof size) How many linear functions f_1, \dots, f_m does Alice need to prove that $\text{Max-Cut}(G) \leq k$? In particular, can we always assume that the proof is polynomial-size, i.e. $m \leq |V|^{O(1)}$? (And, more subtly, that each f_i has bit complexity $|V|^{O(1)}$?)
3. (Algorithm) Do quadratic proofs like this correspond to a useful polynomial-time algorithm?

Before we answer these questions, we'll do some examples to gain a bit of intuition about quadratic proofs.

4.1 Quadratic proofs at least as good as linear ones

To make sure we haven't accidentally made our lives worse, we should make sure that we can find a quadratic proof that $\text{Max-Cut}(G) \leq |E|$. In other words, can we find a sum-of-squares representation of the function


$$|E| - \sum_{(u,v) \in E} \frac{1}{2} - \frac{x_u x_v}{2} = \frac{|E|}{2} + \sum_{(u,v) \in E} \frac{x_u x_v}{2} ?$$

We can do this by observing that $(x_u + x_v)^2 = x_u^2 + x_v^2 + 2x_u x_v = 2 + 2x_u x_v$, so

$$\frac{|E|}{2} + \sum_{(u,v) \in E} \frac{x_u x_v}{2} = \sum_{(u,v) \in E} \frac{1}{4} (x_u + x_v)^2.$$

4.2 Quadratic proofs can be better than linear ones: the triangle

The simplest graph $G = (V, E)$ whose maximum cut is strictly smaller than $|E|$ is the triangle. (Every one and two edge graph is bipartite.) So, let $V = \{1, 2, 3\}$ and edges $E = \{(1, 2), (2, 3), (3, 1)\}$:



The actual Max-Cut value is 2. The best upper bound we could hope for from a linear proof system is $|E| = 3$. Can we find a quadratic proof that $\text{Max-Cut}(G) \leq 2.9 < 3$? In fact, there's a simple way to do so: note that

$$(x_1 + x_2 + x_3)^2 = x_1^2 + x_2^2 + x_3^2 + 2(x_1 x_2 + x_2 x_3 + x_3 x_1) = 3 + 2(x_1 x_2 + x_2 x_3 + x_3 x_1).$$

So,

$$2.25 - \left(\frac{1}{2} - \frac{x_1 x_2}{2}\right) - \left(\frac{1}{2} - \frac{x_2 x_3}{2}\right) - \left(\frac{1}{2} - \frac{x_3 x_1}{2}\right) = \frac{1}{4} (x_1 + x_2 + x_3)^2.$$

This is our first proof which beats the linear proof system!

4.3 Quadratic proofs for random graphs

The triangle graph is pretty simple. Now let's investigate a much larger family of graphs: random graphs from the Erdős-Rényi model $G(n, 1/2)$.

Theorem 4. *With high probability over the choice of $G \sim G(n, 1/2)$, there is a polynomial-size quadratic proof that $\text{Max-Cut}(G) \leq (0.5 + o(1))|E|$.*

To prove Theorem 4, we need a linear algebra refresher.

Positive-Semidefinite Matrices and Cholesky Factorization: For a symmetric matrix $M \in \mathbb{R}^{n \times n}$, the following are equivalent. If M satisfies any/all of them, we say it is *positive semi-definite*, or *PSD*.

1. M has all nonnegative eigenvalues.
2. For all $x \in \mathbb{R}^n$, $\langle x, Mx \rangle \geq 0$.
3. There is a matrix $L \in \mathbb{R}^{n \times n}$ such that $M = LL^\top$. This factorization is called the “Cholesky decomposition” of M .

Operator Norm: For a symmetric matrix $M \in \mathbb{R}^{n \times n}$, we write $\|M\|$ for its operator norm, defined as $\|M\| = \max_{x \in \mathbb{R}^n: \|x\|=1} x^\top Mx = \max |\lambda_i(M)|$ where $\lambda_i(M)$ are the eigenvalues of M .

The key fact we will use about random graphs is that they are so-called “two-sided” spectral expander graphs. “Two-sided” here refers to the fact that both the positive and negative eigenvalues of A are bounded in absolute value. We defer the proof of the following lemma to later in the course, when we will be better equipped to handle random matrices.

Lemma 5. *With high probability over the choice of $G \sim G(n, 1/2)$, the adjacency matrix A of G has the property $\|A - \frac{1}{2}J\| \leq o(n)$, where J denotes the all-ones matrix.*

But, what does the operator norm of a matrix have to do with quadratic proofs?

Lemma 6. *Let $M \in \mathbb{R}^{n \times n}$ be a symmetric matrix. There exist n linear functions f_1, \dots, f_n on n variables such that*

$$\|M\| \cdot \|x\|^2 - x^\top Mx = \sum_{i=1}^n f_i(x)^2.$$

(Here the equality holds as degree-2 polynomials over \mathbb{R}^n .)

Proof of Lemma 6. The matrix $\|M\| \cdot I - M$ is positive semi-definite (where I is the identity matrix), so it has a Cholesky decomposition L . Take $f_i(x) = \sum_{j=1}^n L_{ij}x_j$. \square

Proof of Theorem 4. By Lemma 5, it will suffice to show that the conclusion of Theorem 4 holds under the assumption that the graph G satisfies $\|A - \frac{1}{2}J\| \leq o(n)$. Let $M = A - \frac{1}{2}J$.

We can write the Max-Cut objective function as

$$\sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} = \frac{|E|}{2} - \frac{1}{4} x^\top A x = \frac{|E|}{2} - \frac{1}{4} x^\top J x - \frac{1}{4} x^\top M x = \frac{|E|}{2} - \frac{1}{4} \left(\sum_{i \in V} x_i \right)^2 - \frac{1}{4} x^\top M x,$$

where the $1/4$ shows up because each edge appears twice in $x^\top Ax$, as A is a symmetric matrix. Now we can apply Lemma 6 to $-M$ to find $|V|$ linear functions $f_1, \dots, f_{|V|}$ such that $-\|A - \frac{1}{2}J\| \cdot \|x\|^2 + \sum_{i \in V} f_i(x)^2 = x^\top Mx$. Putting this together with the above, we get

$$\begin{aligned} \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} &= \frac{|E|}{2} - \frac{1}{4} \left(\sum_{i \in V} x_i \right)^2 + \|A - \frac{1}{2}J\| \cdot \|x\|^2 - \sum_{i \in V} f_i(x)^2 \\ &= \frac{|E|}{2} + o(|V|^2) - \frac{1}{4} \left(\sum_{i \in V} x_i \right)^2 - \sum_{i \in V} f_i(x)^2 \\ &= (1 + o(1)) \cdot \frac{|E|}{2} - \frac{1}{4} \left(\sum_{i \in V} x_i \right)^2 - \sum_{i \in V} f_i(x)^2, \end{aligned}$$

where we have used that $\|x\|^2 = |V|$ for every $x \in \{-1, 1\}^V$, and that with high probability $|E| = \Omega(|V|^2)$.

Now we are basically done – subtracting both sides of the above from $(1 + o(1)) \cdot \frac{|E|}{2}$ gives

$$(1 + o(1)) \cdot \frac{|E|}{2} - \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} = \frac{1}{4} \left(\sum_{i \in V} x_i \right)^2 + \sum_{i \in V} f_i(x)^2. \quad \square$$

Exercise 7. Strengthen Theorem 4 to $p = o(1)$. Can you prove it for $p = (\log n)^{O(1)}/n$? How about $p = \Theta(1/n)$?

5 From Quadratic Proofs to an Algorithm for Max-Cut

We have a promising proof system on our hands. How could we use it to design an (approximation) algorithm for Max-Cut?

5.1 Convexity, Duality and Pseudoexpectations

In the Min-Cut setting, where we had a linear proof system, the proof system and the Min-Cut LP were related by LP duality. We would like to find an appropriate generalization of LP duality which will enable us to convert our quadratic proof system into an approximation algorithm for Max-Cut. As a first step, we observe that the set of “possible right-hand sides” for a quadratic proof is convex.

Lemma 8 (Convexity of Sum-of-Squares Polynomials). *For any $d \in \mathbb{N}$, the set of functions $p(x) : \{-1, 1\}^n \rightarrow \mathbb{R}$ for which there exist f_1, \dots, f_m of degree $\leq d$ such that $p(x) = \sum_{i=1}^m f_i(x)^2$ for all $x \in \{-1, 1\}^n$ is convex. We call this set SoS_d .*


Proof. The proof is straightforward: if $p(x) = \sum_{i=1}^m f_i(x)^2$ and $q(x) = \sum_{i=1}^m g_i(x)^2$, then for any $\lambda \in [0, 1]$ we have

$$\lambda p(x) + (1 - \lambda)q(x) = \sum_{i=1}^m (\sqrt{\lambda} f_i(x))^2 + (\sqrt{1 - \lambda} g_i(x))^2. \quad \square$$

So we can phrase the question “is there a quadratic proof that $\text{Max-Cut}(G) \leq k$?” as a convex membership problem: is $k - \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2}$ in SoS_1 ?

Convex Duality 101 (Informal): Recall that $C \subseteq \mathbb{R}^n$ is convex if for every $x, y \in C$, also $\frac{x+y}{2} \in C$. There is a rich theory of convex geometry – we will discuss a tiny portion of it, on the topic of convex duality, and not be particularly formal about it for now. Formal treatments, handling technical issues I won’t discuss here, are available in many textbooks.

Every convex C set looks kind of like this:



The dashed lines are “supporting hyperplanes” of C – they are hyperplanes which lie on one side of C , and intersect (the closure of) C at one or more points. Every point on the boundary of a convex set is intersected by some supporting hyperplane. A “separating” hyperplane lies on one side of the convex set, and need not C at any point. The set of all separating hyperplanes is called the “dual” of C .

Formally, if $C \subseteq \mathbb{R}^n$ is a convex set, a separating hyperplane is an affine function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ such that $f(x) \geq 0$ for all $x \in C$. The hyperplane f is supporting if there is some x in the closure of C with $f(x) = 0$. Convexity of the set of separating hyperplanes is clear: if $f(x), g(x) \geq 0$ for all $x \in C$, then also $\frac{f(x)+g(x)}{2} \geq 0$.

Furthermore, any $x \notin \text{closure}(C)$ if and only if there is a separating hyperplane for x ; that is a separating hyperplane f of C such that $f(x) < 0$.

Duals to Quadratic Proofs Now that we know about convex duality, let’s figure out what the convex dual of SoS_1 is. Let us think of SoS_1 as a subset of $\mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m$, the vector space of multilinear polynomials in variables x_1, \dots, x_n with real coefficients and degree at most 2.² (The “ m ” stands for multilinear – it is not standard notation.) A separating hyperplane for SoS_1 is an affine function $\mathcal{L} : \mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m \rightarrow \mathbb{R}$ such that $\mathcal{L}[p] \geq 0$ for all $p \in \text{SoS}_1$.³ Because of some special structure of SoS_1 (it is actually a convex cone), we’ll be able to assume that \mathcal{L} is actually linear rather than affine, and satisfies a normalization condition. This is captured by the following theorem.

Theorem 9. *Let $p \in \mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m$ and let $\varepsilon > 0$. If $p + \varepsilon \notin \text{SoS}_1$, then there is a linear function $\tilde{\mathcal{L}} : \mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m \rightarrow \mathbb{R}$ satisfying:*

²If we want to be picky about thinking of convex sets as living in \mathbb{R}^N for some value of N , we could identify $\mathbb{R}[x_1, \dots, x_n]_{\leq 2}$ with $\mathbb{R}^{\binom{N}{2}+N+1}$ by identifying a polynomial with the list its of coefficients in the monomial basis.

³We use the notation $\mathcal{L}[p]$ for the function \mathcal{L} applied to the polynomial p , rather than $\mathcal{L}(p)$, for reasons which will be clear later.

1. *Positivity on SoS polynomials:* $\widetilde{\mathbb{E}}[q(x)] \geq 0$ for every $q \in \text{SoS}_1$

2. *Normalization:* $\widetilde{\mathbb{E}}[1] = 1$

3. *Separation:* $\widetilde{\mathbb{E}}[p(x)] < 0$.

Conversely, if such a $\widetilde{\mathbb{E}}$ exists, then $p \notin \text{SoS}_1$.

Proof. The “if” direction is clear, so we prove the “only if”. By the separating hyperplane theorem for convex sets, $p + \varepsilon \notin \text{SoS}_1$ implies that there is a hyperplane \mathcal{L} separating $p + \varepsilon$ from SoS_1 . We claim that the linear function defined by $\mathcal{L}'[q] = \mathcal{L}[q] - \mathcal{L}[0]$ for all q also separates $p + \varepsilon$ from SoS_1 .

We claim that $\mathcal{L}[0] \geq 0$, as otherwise there would be a positive $c > 0$ such that $\mathcal{L}[c] < 0$, by continuity, but since $c = (\sqrt{c})^2$ this can't happen. So, $\mathcal{L}'[p + \varepsilon] = \mathcal{L}[p + \varepsilon] - \mathcal{L}[0] \leq 0$.

It just remains to verify that $\mathcal{L}'[q] \geq 0$ for every $q \in \text{SoS}_1$; then we can scale \mathcal{L}' by a nonnegative scalar to make it satisfy the normalization condition. It will be enough to show that $\mathcal{L}[q] \geq \mathcal{L}[0]$ for every $q \in \text{SoS}_1$. Fix $q \in \text{SoS}_1$ and consider the affine function of a single real variable ε given by $t(\varepsilon) = \mathcal{L}[\varepsilon q] = t_1 \varepsilon + t_0$ for some real numbers t_1, t_0 . We claim that $t_1 \geq 0$, as otherwise by taking ε sufficiently large, we could make $\mathcal{L}[\varepsilon q] < 0$, which is impossible. So, $t(\varepsilon)$ is monotone decreasing in ε . But $t(0) = \mathcal{L}[0]$, so we are done.

Let $\widetilde{\mathbb{E}}[q] = \mathcal{L}[q] / \mathcal{L}[1]$, so that $\widetilde{\mathbb{E}}[1] = 1$. Then $\widetilde{\mathbb{E}}[p + \varepsilon] \leq 0$, hence $\widetilde{\mathbb{E}}[p] < 0$. \square

Definition 10 (Pseudoexpectation over the Hypercube). *We call a linear function $\widetilde{\mathbb{E}} : \mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m$ satisfying the positivity and normalization conditions of Theorem 9 a degree 2 pseudoexpectation.*

We often extend a pseudoexpectation $\widetilde{\mathbb{E}}$ to a linear function on all of $\mathbb{R}[x_1, \dots, x_n]_{\leq 2}$ by defining $\widetilde{\mathbb{E}}[x_i^2] = 1$ and extending linearly. Notice that this extension has the property that two polynomials p, q for which $p(x) = q(x)$ for all $x \in \{\pm 1\}^n$ must have $\widetilde{\mathbb{E}}[p(x)] = \widetilde{\mathbb{E}}[q(x)]$. In particular, for every linear function $q(x)$, $\widetilde{\mathbb{E}}[q(x)^2] \geq 0$.

Duality for Max-Cut Quadratic Proofs Theorem 9 has the following consequence for Max-Cut:

Corollary 11. *Let $G = (V, E)$ be a graph. Then*

$$\inf \left\{ k : k - \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \in \text{SoS}_1 \right\} = \sup \widetilde{\mathbb{E}} \left[\sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \right].$$

Proof. If k is such that $k - \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \notin \text{SoS}_1$, then there exists $\widetilde{\mathbb{E}}$ such that

$$\widetilde{\mathbb{E}} \left[k - \varepsilon - \sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \right] < 0,$$

which rearranges to

$$\widetilde{\mathbb{E}} \left[\sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \right] \geq \widetilde{\mathbb{E}}[k - \varepsilon] = k - \varepsilon.$$

Since this holds for every $\varepsilon > 0$, we have the corollary. \square

At this point, our hoped-for outline of an algorithm for Max-Cut is:

1. Given a graph G , find a pseudoexpectation $\tilde{\mathbb{E}}$ maximizing $\tilde{\mathbb{E}}[\sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2}]$.
2. Use $\tilde{\mathbb{E}}$ to find a cut in G which cuts many edges.

For (1), we can be optimistic, because the set of pseudoexpectations is convex and only $\approx n^2$ -dimensional, when G is an n -vertex graph, and we are trying to maximize a linear function over it. The set of pseudoexpectations isn't a polytope (that is, the feasible region of a linear program), so we can't use linear programming for (1), but we will be able to use a generalization of linear programming called *semidefinite programming*. For now, we will pretend that we can solve (1) in polynomial time, and move on to (2), for which we need to understand better how to think of a pseudoexpectation.

What is a Pseudoexpectation, Anyway? We have a formal definition of a pseudoexpectation, but it's not very intuitive. A good way to think of a degree-2 pseudoexpectation is that it tells us the degree-two moments (that is, means and covariances) of a "fictional" probability distribution over $\{\pm 1\}^n$.

Exercise 12. Let μ be a probability distribution over $\{\pm 1\}^n$ and define $\tilde{\mathbb{E}}[p(x)] = \mathbb{E}_{x \sim \mu} p(x)$. Show that $\tilde{\mathbb{E}}$ is a degree-2 pseudoexpectation.

So, every distribution over the hypercube gives a pseudoexpectation. The reverse is not true.

Exercise 13. Find a degree-2 pseudoexpectation $\tilde{\mathbb{E}}$ over $\{\pm 1\}^3$ which is not induced by any distribution over $\{\pm 1\}^3$.

Even so, it helps to think of a pseudoexpectation as specifying moments of a probability distribution. In the case of Max-Cut, the pseudoexpectation appears to be the second moments of a probability distribution over cuts in G which cut a lot of edges. Our goal is to use these second moments to find an actual cut which cuts a lot of edges.

It is tempting to try and sample from a distribution which induces the pseudoexpectation, but this isn't allowed – there may be no such distribution!

5.2 Gaussian Rounding

Goemans and Williamson had the following amazing idea for how to use a pseudoexpectation to find a cut which cuts a lot of edges. To express it, we need the following lemma.

Lemma 14. Let $\tilde{\mathbb{E}}$ be a degree-2 pseudoexpectation. There is a Gaussian distribution \mathcal{N} over \mathbb{R}^n such that for every $p \in \mathbb{R}[x_1, \dots, x_n]_{\leq 2}^m$, we have $\tilde{\mathbb{E}}[p(x)] = \mathbb{E}_{y \sim \mathcal{N}} p(y)$.

Proof. We let $\mu = \tilde{\mathbb{E}}[x]$ and $\Sigma = \tilde{\mathbb{E}}[(x - \mu)(x - \mu)^\top]$. We claim that there is a Gaussian distribution whose mean and covariance are μ and Σ . We first establish that the matrix Σ is PSD. Let $v \in \mathbb{R}^n$; we need to show that $\langle v, \Sigma v \rangle \geq 0$. But $\langle v, \Sigma v \rangle = \tilde{\mathbb{E}}[\langle v, x - \mu \rangle^2] \geq 0$. Since Σ is PSD, there is a matrix B such that $\Sigma = BB^\top$. Let the Gaussian random vector y be given by sampling $g \sim \mathcal{N}(0, I_{n \times n})$, and letting

$$y = \mu + Bg.$$

□

Goemans-Williamson Rounding Given a degree-2 pseudoexpectation $\widetilde{\mathbb{E}}$, define $\widetilde{\mathbb{E}}'[p(x)] = \frac{1}{2} \widetilde{\mathbb{E}}[p(x)] + \frac{1}{2} \widetilde{\mathbb{E}}[p(-x)]$, so that $\widetilde{\mathbb{E}}'[x] = 0$. Sample y from the Gaussian distribution \mathcal{N} whose mean and covariance are given by $\widetilde{\mathbb{E}}'$. Then, output the cut $S = \{i : y_i \geq 0\}$.

Theorem 15. *Let $G = (V, E)$ be a graph and let $\widetilde{\mathbb{E}}$ be a degree-2 pseudoexpectation. Then the cut S output by Goemans-Williamson rounding satisfies*

$$\mathbb{E} |E(S, \bar{S})| \geq 0.878 \cdot \widetilde{\mathbb{E}} \left[\sum_{(i,j) \in E} \frac{1}{2} - \frac{x_i x_j}{2} \right].$$

Acknowledgements

This lecture is heavily inspired by several lectures in Ryan O'Donnell's Theory's Toolkit course at CMU, as captured in https://youtube.com/playlist?list=PLm3J0oaFux3ZYpFLwwrlv_EHH9wtH6pnX&si=Yb4BSZ4iz75hn2rH.